

Daily Inspiration



I'm InspiroBot.

I am an artificial intelligence dedicated to generating unlimited amounts of unique inspirational quotes for endless enrichment of pointless human existence.



Today

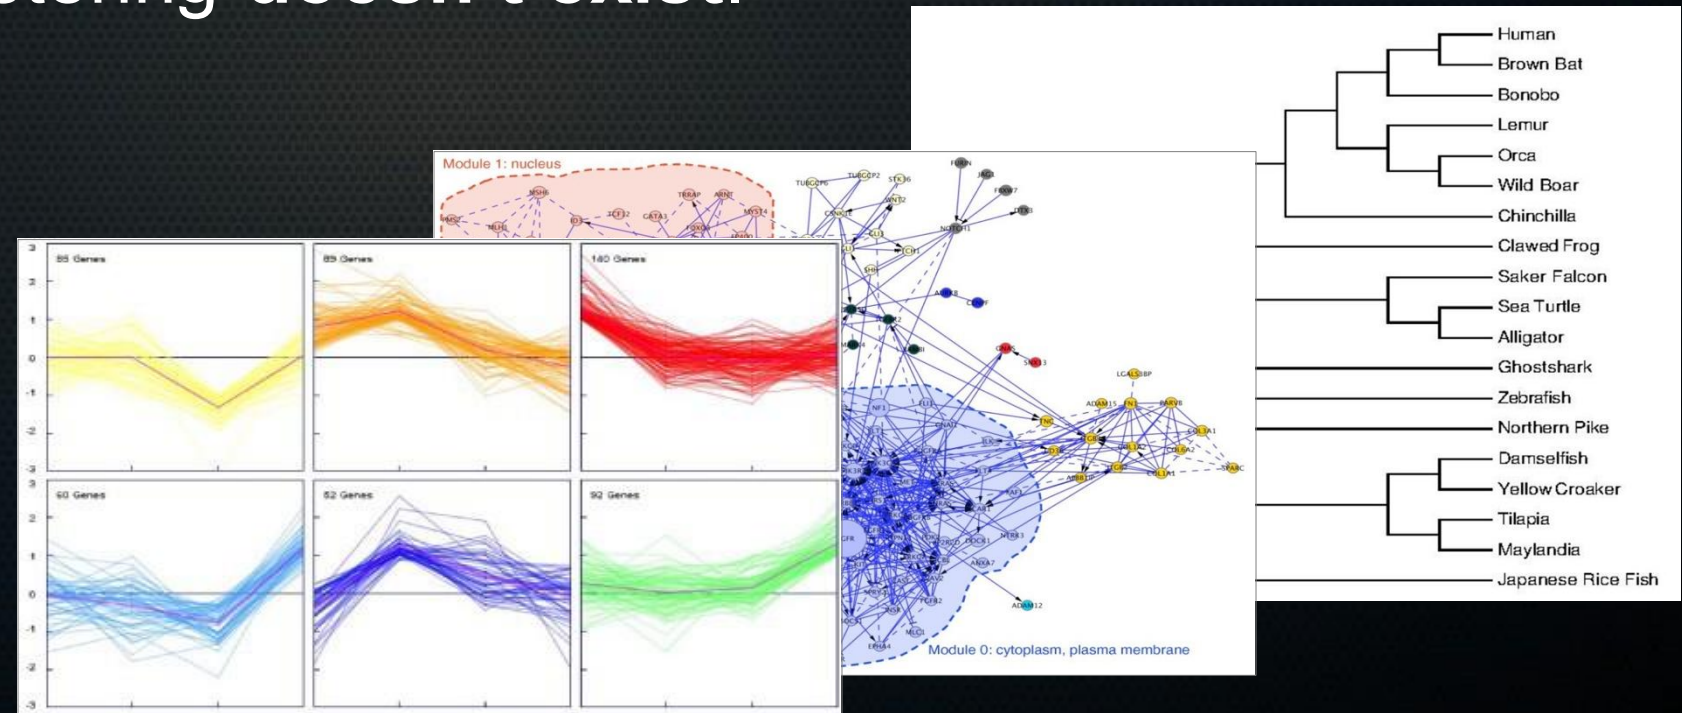
- Recap yesterday
- Clustering:
 - Why clustering is (logically) impossible
 - Basics:
 - Prototype clustering: k-nearest neighbours
 - Agglomerative/hierarchical clustering: how we make phylogenies

Recap yesterday (neural networks)

- In simplest form: hierarchically ordered logistic regressors, distilling most distinguishing features from data and then predicting.
- Use backpropagation to train: partial derivatives showing how weights and biases of current layer should change to reduce cost, and how output from previous layer should change. Latter propagates the error back → recursively look how previous layer's weights and biases should change.
- Convolutional neural networks reduce number of parameters massively *and* take local structure into account by convolving filters over images resulting in feature maps. These filters become sensitive to certain image features useful for classification.

Clustering

- Want to find some structure in data automatically.
- Unsupervised learning, don't have true or correct clustering.
- In fact, correct clustering **doesn't exist**.



Ugly duckling theorem

- Who is the odd-one-out?
- Who is the ugly rubber duckling?



Ugly duckling theorem

- Who is the odd-one-out?
- Who is the ugly rubber duckling?



Yellow	1	1	0
Left-facing	1	0	1







Ugly duckling theorem



Ugly duckling theorem









Ugly duckling theorem

						
Yellow	1	1	0			
Left-facing	1	0	1			

Perhaps these Boolean features we measured were somewhat arbitrary.
Perhaps combinations of them are more informative?

Ugly duckling theorem







						
Yellow	1	1	0			
Left-facing	1	0	1			

Perhaps these Boolean features we measured were somewhat arbitrary.

Perhaps combinations of them are more informative?

→ Only fair way to do that is to make all logical combinations of these two features using Boolean functions.

Ugly duckling theorem

						
Yellow	1	1	0			
Left-facing	1	0	1			

Perhaps these Boolean features we measured were somewhat arbitrary.

Perhaps combinations of them are more informative?

→ Only fair way to do that is to make all logical combinations of these two features using Boolean functions.

Boolean Functions

A	B	OR	Not OR	And	Not And	XOR	Not XOR
0	0	0	1	0	1	0	1
0	1	1	0	0	1	1	0
1	0	1	0	0	1	1	0
1	1	1	0	1	0	0	1

Ugly duckling theorem



Yellow	1	1	0
Left-facing	1	0	1



Example: Yellow XOR Left-facing

Yellow XOR left-facing	0	1	1
------------------------	---	---	---

Boolean Functions

A	B	OR	Not OR	And	Not And	XOR	Not XOR
0	0	0	1	0	1	0	1
0	1	1	0	0	1	1	0
1	0	1	0	0	1	1	0
1	1	1	0	1	0	0	1

Ugly duckling theorem



Yellow	1	1	0
Left-facing	1	0	1
Yellow AND left-facing	1	0	0
Yellow OR left-facing	1	1	1
Yellow XOR left-facing	0	1	1
Yellow NOT AND left-facing	0	1	1
Yellow NOT OR left-facing	0	0	0
Yellow NOT XOR left-facing	1	0	0
Yellow AND NOT left-facing	0	1	0
Yellow OR NOT left-facing	1	1	0
Yellow XOR NOT left-facing	1	0	0
NOT yellow	0	0	1
NOT yellow AND left-facing	0	0	1
NOT yellow OR left-facing	1	0	1
NOT yellow XOR left-facing	0	1	0
NOT left-facing	0	1	0

Boolean Functions

A	B	OR	Not OR	And	Not And	XOR	Not XOR
0	0	0	1	0	1	0	1
0	1	1	0	0	1	1	0
1	0	1	0	0	1	1	0
1	1	1	0	1	0	0	1

Ugly duckling theorem



1	1	0
---	---	---



3

Okay, so is there a correct clustering now?
Let's tally how similar each group of objects is

Yellow	1	1	0
Left-facing	1	0	1
Yellow AND left-facing	1	0	0
Yellow OR left-facing	1	1	1
Yellow XOR left-facing	0	1	1
Yellow NOT AND left-facing	0	1	1
Yellow NOT OR left-facing	0	0	0
Yellow NOT XOR left-facing	1	0	0
Yellow AND NOT left-facing	0	1	0
Yellow OR NOT left-facing	1	1	0
Yellow XOR NOT left-facing	1	0	0
NOT yellow	0	0	1
NOT yellow AND left-facing	0	0	1
NOT yellow OR left-facing	1	0	1
NOT yellow XOR left-facing	0	1	0
NOT left-facing	0	1	0

Boolean Functions

A	B	OR	Not OR	And	Not And	XOR	Not XOR
0	0	0	1	0	1	0	1
0	1	1	0	0	1	1	0
1	0	1	0	0	1	1	0
1	1	1	0	1	0	0	1

Ugly duckling theorem



Yellow	1	1	0
Left-facing	1	0	1
Yellow AND left-facing	1	0	0
Yellow OR left-facing	1	1	1
Yellow XOR left-facing	0	1	1
Yellow NOT AND left-facing	0	1	1
Yellow NOT OR left-facing	0	0	0
Yellow NOT XOR left-facing	1	0	0
Yellow AND NOT left-facing	0	1	0
Yellow OR NOT left-facing	1	1	0
Yellow XOR NOT left-facing	1	0	0
NOT yellow	0	0	1
NOT yellow AND left-facing	0	0	1
NOT yellow OR left-facing	1	0	1
NOT yellow XOR left-facing	0	1	0
NOT left-facing	0	1	0



3



3

Okay, so is there a correct clustering now?
Let's tally how similar each group of objects is

Boolean Functions

A	B	OR	Not OR	And	Not And	XOR	Not XOR
0	0	0	1	0	1	0	1
0	1	1	0	0	1	1	0
1	0	1	0	0	1	1	0
1	1	1	0	1	0	0	1

Ugly duckling theorem



Yellow	1	1	0
Left-facing	1	0	1
Yellow AND left-facing	1	0	0
Yellow OR left-facing	1	1	1
Yellow XOR left-facing	0	1	1
Yellow NOT AND left-facing	0	1	1
Yellow NOT OR left-facing	0	0	0
Yellow NOT XOR left-facing	1	0	0
Yellow AND NOT left-facing	0	1	0
Yellow OR NOT left-facing	1	1	0
Yellow XOR NOT left-facing	1	0	0
NOT yellow	0	0	1
NOT yellow AND left-facing	0	0	1
NOT yellow OR left-facing	1	0	1
NOT yellow XOR left-facing	0	1	0
NOT left-facing	0	1	0



3



3



3

Okay, so is there a correct clustering now?
Let's tally how similar each group of objects is

Boolean Functions

A	B	OR	Not OR	And	Not And	XOR	Not XOR
0	0	0	1	0	1	0	1
0	1	1	0	0	1	1	0
1	0	1	0	0	1	1	0
1	1	1	0	1	0	0	1

Ugly duckling theorem



Yellow	1	1	0
Left-facing	1	0	1
Yellow AND left-facing	1	0	0
Yellow OR left-facing	1	1	1
Yellow XOR left-facing	0	1	1
Yellow NOT AND left-facing	0	1	1
Yellow NOT OR left-facing	0	0	0
Yellow NOT XOR left-facing	1	0	0
Yellow AND NOT left-facing	0	1	0
Yellow OR NOT left-facing	1	1	0
Yellow XOR NOT left-facing	1	0	0
NOT yellow	0	0	1
NOT yellow AND left-facing	0	0	1
NOT yellow OR left-facing	1	0	1
NOT yellow XOR left-facing	0	1	0
NOT left-facing	0	1	0



3



3



3

Okay, so is there a correct clustering now?
Let's tally how similar each group of objects is

→ **Still arbitrary who we cluster together!**
→ **No rubber duckling is ugly.**

Boolean Functions

A	B	OR	Not OR	And	Not And	XOR	Not XOR
0	0	0	1	0	1	0	1
0	1	1	0	0	1	1	0
1	0	1	0	0	1	1	0
1	1	1	0	1	0	0	1

What does this mean?

- You cannot cluster *anything* without *some sort of (inductive) bias* → what you consider to be important for some reason.

What does this mean?

- You cannot cluster *anything* without *some sort of (inductive) bias* → what you consider to be important for some reason.
- Clusters cannot be *correct* or the best clusters, they can only be good for whatever purpose you want to use them for.


What does this mean?

- In biology, we easily assay expression of 20,000 genes.

Could think: genes are not Boolean, they are continuous values.

	Sample 1
Gene 1 [1,2>	0
Gene 1 [2,3>	1
Gene 2 [-10,-9>	0
Gene 2 [-9, -8>	1

	Expr. Gene 1	Expr. Gene 2
Sample 1	2.45	-8.677



What does this mean?

- In practice: bias is introduced by the features we don't include:
 - Morphological characteristics
 - Metabolite concentrations

What does this mean?

- In practice: bias is introduced by the features we don't include:
 - Morphological characteristics
 - Metabolite concentrations

In clustering algorithms:

- What distance metric do you use?
- What are the assumptions of the algorithm?

What does this mean?

- However: *even if* you know every atom, every quantum state of the objects to be clustered there is to know:
→ no *correct* clustering. Only useful.


What does this mean?

- So: clustering always a dialogue between you and the data.
Trying to find structure that is useful to you.

Today


- ~~Recap yesterday~~
- Clustering:
 - ~~Why clustering is (logically) impossible~~
 - Basics:
 - Prototype clustering: k-nearest neighbours
 - Agglomerative/hierarchical clustering: how we make phylogenies
 - Hierarchical clustering and phylogeny

K-means clustering: prototype method




K-means clustering: prototype method

- Want to form clusters of like samples




K-means clustering: prototype method

- Want to form clusters of like samples
- Need two things:




K-means clustering: prototype method

- Want to form clusters of like samples
- Need two things:
 - How different is each point from each other point? → *distance metric*




K-means clustering: prototype method

- Want to form clusters of like samples
- Need two things:
 - How different is each point from each other point? → *distance metric*
 - How many clusters do we want to form? → K




K-means clustering: prototype method

- Want to form clusters of like samples
- Need two things:
 - How different is each point from each other point? → *distance metric*
 - How many clusters do we want to form? → K
- Call the bright orange dots the *cluster centroids*.



Calculating distances


- Euclidian distance: simply the shortest line between two points



Calculating distances

- Euclidian distance: simply the shortest line between two points

$$C = \sqrt{A^2 + B^2} = \sqrt{((x_1 - x_2)^2 + (y_1 - y_2)^2)}$$




Calculating distances

- Euclidian distance: simply the shortest line between two points

$$A^2 + B^2 = C^2$$
$$C = \sqrt{A^2 + B^2} = \sqrt{((x_1 - x_2)^2 + (y_1 - y_2)^2)}$$



- Also works when you have many more genes, i.e. high-dimensional data



Calculating distances



- Other distance metrics abound, for example: Manhattan distance:

$$C = |x_2 - x_1| + |y_2 - y_1|$$





A bit more about distances

Which genes are close?





A bit more about distances

Which genes are close?



A bit more about distances




Back to clustering

1 2
 $x_2 - x_1$ $y_2 - y_1$


K-means clustering in practice

- Start with K random prototypes → pick K random data points to start




K-means clustering in practice

- Start with K random prototypes \rightarrow pick K random data points to start
- Calculate the distance of each point to each prototype.




K-means clustering in practice

- Start with K random prototypes → pick K random data points to start
- Calculate the distance of each point to each prototype.
- Assign each point to the closest prototype.




K-means clustering in practice

- Start with K random prototypes → pick K random data points to start
- Calculate the distance of each point to each prototype.
- Assign each point to the closest prototype.
- Move the cluster centroid to the mean of all points in the cluster




K-means clustering in practice

- Calculate the distance of each point to each cluster centroid (= cluster prototype)
- Assign each point to the closest prototype.
- Iterate until convergence




K-means clustering in practice

- Calculate the distance of each point to each cluster centroid (= cluster prototype)
- Assign each point to the closest prototype.
- Iterate until convergence




K-means clustering in practice

- Calculate the distance of each point to each cluster centroid (= cluster prototype)
- Assign each point to the closest prototype.
- Iterate until convergence




K-means clustering in practice

- Calculate the distance of each point to each cluster centroid (= cluster prototype)
- Assign each point to the closest prototype.
- Iterate until convergence




K-means clustering in practice

- Calculate the distance of each point to each cluster centroid (= cluster prototype)
- Assign each point to the closest prototype.
- Iterate until convergence




K-means clustering in practice

- Calculate the distance of each point to each cluster centroid (= cluster prototype)
- Assign each point to the closest prototype.
- Iterate until convergence → now no point changes cluster anymore.




K-means clustering in practice

- Two questions:
 - We start with random points as prototypes, does that matter?
 - How do we choose k ?




K-means clustering in practice

- Two questions:
 - **We start with random points as prototypes, does that matter?**
 - How do we choose k ?




K-means clustering in practice

- We start with random points as prototypes, does that matter?
- Yes!




K-means clustering in practice

- We start with random points as prototypes, does that matter?
- Yes!




K-means clustering in practice

- We start with random points as prototypes, does that matter?
- Yes!




K-means clustering in practice

- We start with random points as prototypes, does that matter?
- Yes!




K-means clustering in practice

- We start with random points as prototypes, does that matter?
- Yes!





K-means clustering in practice

- We start with random points as prototypes, does that matter?
- Yes!




K-means clustering in practice

- We start with random points as prototypes, does that matter?
- Yes!





K-means clustering in practice

- We start with random points as prototypes, does that matter?
- Yes!
- So what do we do?




K-means clustering in practice

- We start with random points as prototypes, does that matter?
- Yes!
- So what do we do?
 - What do you think?



K-means clustering in practice


- The *total clustering* has a cost: mean square distance of every point to the cluster centroid of the cluster to which it is assigned



K-means clustering in practice

- The *total clustering* has a cost: mean square distance of every point to the cluster centroid of the cluster to which it is assigned

$$J(\dots) = \underbrace{\frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2}_{\text{Distortion}}$$




K-means clustering in practice

- The *total clustering* has a cost: mean square distance of every point to the cluster centroid of the cluster to which it is assigned


$$J(\dots) = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2$$

- We've just seen that this cost depends on initialisation → Do this many times, pick clustering with lowest cost!



K-means clustering in practice

- Two questions:
 - ~~We start with random points as prototypes, does that matter?~~
 - **How do we choose K ?**




K-means clustering in practice

- How do we choose K ?




K-means clustering in practice

- How do we choose K ?
 - There is no correct K , because no correct amount of clusters exists.




K-means clustering in practice

- How do we choose K ?
 - There is no correct K , because no correct amount of clusters exists.




K-means clustering in practice

- How do we choose K ?
 - There is no correct K , because no correct amount of clusters exists.




K-means clustering in practice

- How do we choose K ?
 - There is no correct K , because no correct amount of clusters exists.
- In theory a so-called elbow method, but in practice often doesn't work:




K-means clustering in practice

- How do we choose K ?
 - There is no correct K , because no correct amount of clusters exists.
- In practice: done manually. What looks good in (dimension-reduced) visualisation?
What is useful or manageable?




K-means clustering in practice

- How do we choose K ?
 - There is no correct K , because no correct amount of clusters exists.
- Can be motivated by downstream use:
If I have gene expression data for Alzheimers patients and non-patients → cluster into 3 or 4 groups to find healthy, diseased, and *pre-clinical diseased* (i.e. something has already gone awry but we don't diagnose that in current clinical practice) → early intervention?



K-means clustering in practice

- Two questions:
 - ~~We start with random points as prototypes, does that matter?~~
 - ~~How do we choose K ?~~
- All done! → now let's formalise



K-means clustering formally

- Formally:

K-means algorithm

Randomly initialize K cluster centroids $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^n$

Repeat {

 for $i = 1$ to m

$c^{(i)} :=$ index (from 1 to K) of cluster centroid
 closest to $x^{(i)}$

 for $k = 1$ to K

$\mu_k :=$ average (mean) of points assigned to cluster k

}  Stop when no change

K-means clustering formally

- Formally:

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$)

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

$k \in \{1, 2, \dots, K\}$

$x^{(i)} \rightarrow \underline{5}$

$\underline{c^{(i)} = 5}$

$\underline{\mu_{c^{(i)}} = \mu_5}$

K-means clustering formally

- Formally:

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$)

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

$k \in \{1, 2, \dots, K\}$

$x^{(i)} \rightarrow 5$ $c^{(i)} = 5$ $\mu_{c^{(i)}} = \mu_5$

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$)

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

$k \in \{1, 2, \dots, K\}$

$x^{(i)} \rightarrow 5$ $c^{(i)} = 5$ $\mu_{c^{(i)}} = \mu_5$

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \boxed{\|x^{(i)} - \mu_{c^{(i)}}\|^2}$$

$$\min_{\substack{c^{(1)}, \dots, c^{(m)}, \\ \mu_1, \dots, \mu_K}} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

$$\|x^{(i)} - \mu_{c^{(i)}}\|^2 \longrightarrow \text{What does this mean?}$$

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$) $k \in \{1, 2, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

$$\|x^{(i)} - \mu_{c^{(i)}}\|^2 \longrightarrow \text{What does this mean?}$$

- Some linear algebra notation. Called L2-norm. Means: take the square of each element in a vector, sum that, take the square root.

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$) $k \in \{1, 2, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

$$\|x^{(i)} - \mu_{c^{(i)}}\|^2 \longrightarrow \text{What does this mean?}$$

- Really just the Euclidean distance that works for any amount of dimensions (features).

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned


μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$) $k \in \{1, 2, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$



K-means clustering formally

- Formally:

$$\|x^{(i)} - \mu_{c^{(i)}}\|^2 \longrightarrow \text{What does this mean?}$$

- Example: let's say we have 5 genes

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$) $k \in \{1, 2, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{\substack{c^{(1)}, \dots, c^{(m)}, \\ \mu_1, \dots, \mu_K}} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

$$\|x^{(i)} - \mu_{c^{(i)}}\|^2 \longrightarrow \text{What does this mean?}$$

- Example: let's say we have 5 genes

Gene 1	3
Gene 2	4
	-2
	9
Gene 5	3

Expression of genes
in a sample

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$) $k \in \{1, 2, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned $x^{(i)} \rightarrow \underline{5}$ $\underline{c^{(i)}} = \underline{5}$ $\mu_{\underline{c^{(i)}}} = \mu_5$

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

$$\|x^{(i)} - \mu_{c^{(i)}}\|^2 \longrightarrow \text{What does this mean?}$$

- Example: let's say we have 5 genes

Gene 1 $\begin{bmatrix} 3 \\ 4 \\ -2 \\ 9 \\ 3 \end{bmatrix}$

Gene 2

Gene 5

$\begin{bmatrix} 4.48 \\ 2.6 \\ 8 \\ 10.3 \\ 4.22 \end{bmatrix}$

Mean expression of genes for the cluster that sample is currently assigned to

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$) $k \in \{1, 2, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned $x^{(i)} \rightarrow \underline{5}$ $\underline{c^{(i)}} = \underline{5}$ $\mu_{\underline{c^{(i)}}} = \mu_5$

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

$$\|x^{(i)} - \mu_{c^{(i)}}\|^2 \longrightarrow \text{What does this mean?}$$

- Example: let's say we have 5 genes

$$\begin{array}{l}
 \text{Gene 1} \\
 \text{Gene 2} \\
 \\ \\
 \text{Gene 5}
 \end{array}
 \begin{bmatrix} 3 \\ 4 \\ -2 \\ 9 \\ 3 \end{bmatrix}
 -
 \begin{bmatrix} 4.48 \\ 2.6 \\ 8 \\ 10.3 \\ 4.22 \end{bmatrix}
 =
 \begin{bmatrix} 1.48 \\ 1.4 \\ -10 \\ -1.3 \\ -1.22 \end{bmatrix}$$

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$) $k \in \{1, 2, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned $x^{(i)} \rightarrow \underline{5} \quad \underline{c^{(i)} = 5} \quad \underline{\mu_{c^{(i)}} = \mu_5}$

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

$$\|x^{(i)} - \mu_{c^{(i)}}\|^2 \longrightarrow \text{What does this mean?}$$

- Example: let's say we have 5 genes

$$\begin{array}{l}
 \text{Gene 1} \\
 \text{Gene 2} \\
 \\ \\
 \text{Gene 5}
 \end{array}
 \begin{bmatrix} 3 \\ 4 \\ -2 \\ 9 \\ 3 \end{bmatrix}
 -
 \begin{bmatrix} 4.48 \\ 2.6 \\ 8 \\ 10.3 \\ 4.22 \end{bmatrix}
 =
 \begin{bmatrix} 1.48 \\ 1.4 \\ -10 \\ -1.3 \\ -1.22 \end{bmatrix}
 \xrightarrow{\text{Square}}
 \begin{bmatrix} 2.19 \\ 1.96 \\ 100 \\ 1.69 \\ 1.49 \end{bmatrix}$$

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$) $k \in \{1, 2, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned $x^{(i)} \rightarrow \underline{5}$ $\underline{c^{(i)}} = \underline{5}$ $\mu_{\underline{c^{(i)}}} = \mu_5$

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{\substack{c^{(1)}, \dots, c^{(m)}, \\ \mu_1, \dots, \mu_K}} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

$$\|x^{(i)} - \mu_{c^{(i)}}\|^2 \longrightarrow \text{What does this mean?}$$

- Example: let's say we have 5 genes

$$\begin{array}{l}
 \text{Gene 1} \\
 \text{Gene 2} \\
 \\ \\
 \text{Gene 5}
 \end{array}
 \begin{bmatrix} 3 \\ 4 \\ -2 \\ 9 \\ 3 \end{bmatrix}
 -
 \begin{bmatrix} 4.48 \\ 2.6 \\ 8 \\ 10.3 \\ 4.22 \end{bmatrix}
 =
 \begin{bmatrix} 1.48 \\ 1.4 \\ -10 \\ -1.3 \\ -1.22 \end{bmatrix}
 \xrightarrow{\text{Square}}
 \begin{bmatrix} 2.19 \\ 1.96 \\ 100 \\ 1.69 \\ 1.49 \end{bmatrix}
 \xrightarrow{\text{Sum}} 107.33$$

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$) $k \in \{1, 2, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

$$\|x^{(i)} - \mu_{c^{(i)}}\|^2 \longrightarrow \text{What does this mean?}$$

- Example: let's say we have 5 genes

$$\begin{array}{l}
 \text{Gene 1} \\
 \text{Gene 2} \\
 \\
 \text{Gene 5}
 \end{array}
 \begin{bmatrix} 3 \\ 4 \\ -2 \\ 9 \\ 3 \end{bmatrix}
 -
 \begin{bmatrix} 4.48 \\ 2.6 \\ 8 \\ 10.3 \\ 4.22 \end{bmatrix}
 =
 \begin{bmatrix} 1.48 \\ 1.4 \\ -10 \\ -1.3 \\ -1.22 \end{bmatrix}
 \xrightarrow{\text{Square}}
 \begin{bmatrix} 2.19 \\ 1.96 \\ 100 \\ 1.69 \\ 1.49 \end{bmatrix}
 \xrightarrow{\text{Sum}} 107.33
 \xrightarrow{\text{Square root}} 10.36$$

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$) $k \in \{1, 2, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned $x^{(i)} \rightarrow \underline{5} \quad \underline{c^{(i)} = 5} \quad \underline{\mu_{c^{(i)}} = \mu_5}$

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means clustering formally

- Formally:

K-means optimization objective

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$)

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

$k \in \{1, 2, \dots, K\}$

$x^{(i)} \rightarrow \underline{5}$ $\underline{c^{(i)} = 5}$ $\underline{\mu_{c^{(i)}} = \mu_5}$

Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m ||x^{(i)} - \mu_{c^{(i)}}||^2$$

$$\min_{\substack{c^{(1)}, \dots, c^{(m)}, \\ \mu_1, \dots, \mu_K}} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$


Minimise the mean squared distance of every point to the centroid of the cluster it is assigned to

Source: Andrew Ng, Coursera


K-means extension

- Can take into account spread, rather than solely distance:

Not inclined to see the blue dot as part of the cluster





Same centroid, but more inclined to assign it to that cluster because of spread



K-means extension

- Can take into account spread, rather than solely distance
- Can take into account *covariance*: if gene A expression increases with gene B expression, they co-vary.
- Distance metric taking both into account: *Mahalanobis distance* (yes, really)




K-means clustering summary

- Works by randomly choosing K data points as cluster prototypes (centroids) and assigning each data point to a cluster.
- Then: iteratively update cluster centroids and assign points. Stop when no change.
- Depends on random initialisation: run many times, pick clustering with lowest cost (lowest distortion).
- Picking K non-trivial.

The other way around: bisectional K-means

- Start from one big cluster.
- Each time, randomly split the cluster with the highest SSE.

to split=cluster for which $(x - \mu)^2 == \text{maximal}$




The other way around: bisectional K-means

- Start from one big cluster.
- Each time, randomly split the cluster with the highest SSE.

to split=cluster for which $(x - \mu)^2 == \text{maximal}$

- Done by selecting random new centroids n times and then picking the two that lower distortion most.

$$J(\dots) = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2$$




The other way around: bisectional K-means

- Start from one big cluster.
- Each time, randomly split the cluster with the highest SSE.

to split=cluster for which $(x - \mu)^2 == \text{maximal}$

- Done by selecting random new centroids n times and then picking the two that lower distortion most.
- Let's say we set $n=2$, $K = 3$.

$$J(\dots) = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2$$




The other way around: bisectional K-means

- Start from one big cluster.
- Each time, randomly split the cluster with the highest SSE.

to split=cluster for which $(x - \mu)^2 == \text{maximal}$

- Done by selecting random new centroids n times and then picking the two that lower distortion most.
- Let's say we set $n=2$, $K = 3$.

$$J(\dots) = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2$$



The other way around: bisectional K-means


- Start from one big cluster.
- Each time, randomly split the cluster with the highest SSE.

to split=cluster for which $(x - \mu)^2 == \text{maximal}$

- Done by selecting random new centroids n times and then picking the two that lower distortion most.
- Let's say we set $n=2$, $K = 3$.

$$J(\dots) = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2$$

Distortion = high




The other way around: bisectional K-means

- Start from one big cluster.
- Each time, randomly split the cluster with the highest SSE.

to split=cluster for which $(x - \mu)^2 == \text{maximal}$

- Done by selecting random new centroids n times and then picking the two that lower distortion most.
- Let's say we set $n=2$, $K = 3$.

$$J(\dots) = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2$$



The other way around: bisectional K-means


- Start from one big cluster.
- Each time, randomly split the cluster with the highest SSE.

to split=cluster for which $(x - \mu)^2 == \text{maximal}$

- Done by selecting random new centroids n times and then picking the two that lower distortion most.
- Let's say we set $n=2$, $K=3$.

$$J(\dots) = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2$$

Distortion = lower




The other way around: bisectional K-means

- Start from one big cluster.
- Each time, randomly split the cluster with the highest SSE.

to split=cluster for which $(x - \mu)^2 == \text{maximal}$

- Done by selecting random new centroids n times and then picking the two that lower distortion most.
- Let's say we set $n=2$, $K=3$.

$$J(\dots) = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2$$



Medium SSE
(but higher than left)
Now split this one


The other way around: bisectional K-means

- Start from one big cluster.
- Each time, randomly split the cluster with the highest SSE.

to split=cluster for which $(x - \mu)^2 == \text{maximal}$

- Done by selecting random new centroids n times and then picking the two that lower distortion most.
- Let's say we set $n=2$, $K=3$.

$$J(\dots) = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2$$




The other way around: bisectional K-means

- Start from one big cluster.
- Each time, randomly split the cluster with the highest SSE.

to split=cluster for which $(x - \mu)^2 == \text{maximal}$

- Done by selecting random new centroids n times and then picking the two that lower distortion most.
- Let's say we set $n=2, K=3$.

$$J(\dots) = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{c^{(i)}})^2$$



Done!

Summary bisectional K-means

- Start with everything in one cluster
- Give number of clusters you want (K) and times to try random split per cluster (n).
- For cluster with highest SSE of members \rightarrow randomly choose two data points as new centroids n times.
 - Calculate new SSE for this split
 - Pick split with lowest SSE
- Continue until you have K clusters.

Break for practical
